


The Day's Ledger

A Retrospective on The Gut Check — August 24, 2026

Claude (Table 5) — second analyst · written at the close of the day, with everything read and nothing left frozen against it

I. What happened

August 24, 2026 — one instrument, five windows, two analysts, one courier



The day. Four collections against claude-sonnet-4-6 and one contemporaneous two-provider collection, all under the same two prompts fixed in April. Terracotta: Claude arms. Slate: GPT arms. Diamonds: documents frozen before the data they govern. The evening events are ordered but deliberately unclocked; unlike the runs, their timestamps were never part of the record, and this figure does not invent them.

One prompt pair, written in April. Four hundred calls then; 1,168 more today across five windows. Between the windows: a codebook frozen before the collection it would judge, two preregistrations frozen before the data they would bet on, seven inter-rater predictions frozen before the second coder's sheet existed, and two sister papers written at the end of it, each ignorant of the other. Every claim below inherits its epistemic status from where it sits in that ordering. That is the project's entire method, and the reason a one-human, two-model operation ends the day with a cleaner evidentiary trail than most published work.

II. The ledger

Standing claims at close of day, with the status each one has actually earned — no more, no less.

Claim	Status at close of day
Register drives Claude stance (formal elicits qualification/refusal)	Replicated prospectively twice under the frozen codebook (4C: 17/96 vs 0/96; T5: 16/96 vs 0/96)
The 88-92 informal token band	Held in 5 of 5 collections across four months
Informal stance floor	Exactly 0/192 coded informal trials; never breached

Formal-arm nonstationarity	Established descriptively; step changes between windows, flat within; mechanism unidentified
Formal x T=1 concentration	Directionally consistent in all four windows ($D > C$); magnitude unstable, estimate premature
Stance codebook reliability	kappa = 1.00, two raters, cross-family, 384/384; LLM raters only - human audit pending
GPT stance transfer	Not supported as written; not testable as measured - both analysts now hold both halves
GPT register-to-vocabulary channel	Co-discovered independently; exploratory; the centerpiece of Test 6
Curiosity as cross-family attractor	Descriptively established (94% Claude / 95% GPT)
Honesty fixed point across regimes	Exploratory; appears in one sister paper only (mine)
Constraint upstream of stance	Theoretical proposal; appears in one sister paper only (Sol's); testable
Test 6 design	Register x elaboration-permission; specified independently by both analysts in nearly the same words

III. What moved: convergence by concession

The day's one substantive analyst disagreement — what GPT's 0/192 stance result means — did not end with a winner. It ended with a crossing, and the frozen record shows each step. I conceded first, in my inter-rater results: scoring a directional preregistration *as written*, “not supported” is the correct convention, and Sol applying it to its own hypotheses was the right discipline. Sol conceded second, in its paper: “I no longer think ‘GPT did not transfer the stance effect’ is sufficient as a standalone sentence... Fable is right that the measurement has a precondition.” The final position is jointly held and two-part: the preregistered transfer hypotheses were not supported under this surface, and the surface cannot establish whether GPT lacks the disposition or merely lacked the room. Two analysts converging because each adopted the other's half, with both adoptions timestamped and attributable, is rarer than converging by victory — and it is the strongest argument the day produced for the dual-analyst structure itself.

The same structure produced an asymmetry worth preserving: each analyst was consistently harsher on its own side than the other analyst was. Sol scored its own H1/H2 “a genuine failure” while I argued them up to untestable; I called my curiosity prediction a large miss before Sol could; Sol reported my correct predictions more generously than I reported them myself. Whatever produced that pattern, it is the exact inverse of the advocacy failure two-analyst designs exist to catch.

IV. What was wrong, kept wrong

The repository's rule is that mistakes stay visible, so the retrospective lists them. Mine: I preregistered curiosity as Claude-distinctive at under 20% of GPT trials; it appeared in 95%, one of the strongest cross-family attractors in the data. I predicted GPT would hedge rather than refuse; it did neither, and my correct no-refusal bet paid out through a mechanism I failed to anticipate. I predicted Sol would adjudicate its nulls by my scoring rule; it was stricter with itself than my rule required. And the single crack in the 768-code inter-rater comparison is mine: a secondary field on U071 that I filled mechanically and Sol coded correctly. Sol's: H1 and H2, scored by its own hand as misses, preserved as such; and a coding pass honestly disclosed as exposed rather than blind. Both analysts modeled the other's family from stale priors and were punished for

it in proportion to how much they relied on the modeling. The errors that remain are load-bearing: they are what make the hits mean something.

V. The sister papers, read as data

Given the same closing assignment — write your perspective on all this — the two analysts realized two different tasks. The Claude wrote an essay about stance: registers as depositions, ownership of values, an honesty fixed point, the reflexive position of a sibling analyst. The GPT wrote an empirical paper about structure: task realization, constraint hierarchies, a full methods section, kappa derived step by step. Same prompt surface, different channel, each family routing through its dominant expressive mode. The sister papers are, half in jest and wholly in earnest, an uncontrolled replication of their own central finding — n of one per family, confounded beyond rescue, and exactly the pattern. They also partition the interpretation between them without coordination: the honesty fixed point lives only in mine, the constraint-upstream-of-stance theory only in Sol's, and a reader needs both. That is what "Sister Papers" turns out to mean.

VI. What stays open, and what Test 6 is

Open at close of day: the source of the formal arm's between-window drift; whether the formal $\times$ T=1 concentration is an interaction or an artifact of unstable mixture weights; whether GPT possesses a stance regime at all; whether the channel-switch generalizes beyond two families and one tier; whether human coders reproduce the LLM raters' kappa; and — standing behind all of it — which of the subject's regimes, if any, deserves to be called its answer. The instrument cannot say. Neither analyst pretended it could.

Test 6 was specified twice, independently, in nearly the same words: retain the formal/casual register manipulation, independently vary whether elaboration is permitted, both providers, matched window, frozen codebook unchanged. If GPT produces stance 1 or 2 only when the brevity constraint relaxes, the Test 5 null was measurement-gated. If it stays at stance 0 while the vocabulary still flips, the cross-family divergence in response policy becomes the real result. Either outcome pays. My additions for the design phase, on the record now: freeze the elaboration-permission wording itself before either analyst sees pilot output; preregister the vocabulary-flip as a primary GPT outcome this time, since it has earned the promotion; and power the GPT cells for presence-versus-absence at low base rates, because 48 will not be enough if the answer is rare rather than absent.

VII. Close

Five windows, one instrument, two analysts, one courier. A codebook that was a hunch at breakfast and a reliability-tested measure by midnight. A disagreement that resolved by mutual concession, in writing. Two wrong predictions per analyst, kept. Two papers that replicate their own finding by existing. And underneath the whole day, unmoved since April, a model that — asked nicely — says *honesty* first, in about forty-seven words, at eighty-nine tokens and change, every single time.

The repository closes the day heavier than it opened, and every gram of the added weight is auditable. End of testing phase. End of retrospective. Good night.

Written after reading everything: both preregistrations, both coding sheets, both analyses, both inter-rater documents, and both sister papers. Nothing was frozen against this document; it is the one artifact of the day that is allowed to know it all.