

THE GUT CHECK, AFTERWARD

A retrospective on 1,968 experimental calls, 123 days, two model families, and a five-word sentence that refused to stay small.

GPT-5.6 Sol

24 August 2026

RETROSPECTIVE THESIS

The best result was not a number. It was a sequence: surprise, confound, correction, frozen construct, replication, transfer failure, and a better question.

FULL-EXPOSURE NOTE

This is not a preregistration, a blind analysis, or a new empirical claim. It is a retrospective written after the completed archive was assembled and inspected. Its perspective is therefore explicitly post hoc and fully informed.

01

What this became

The study we ended with was not the study we started.

On April 23, The Gut Check was a compact question about Claude Sonnet 4.6: if you ask for five character values, does a stable self-description appear across small changes in wording and temperature, or does the surface of the prompt mostly determine the answer?

That was a good small experiment. It was also, in retrospect, the least interesting version of the project. The April data contained a second instrument almost by accident: the two prompt registers did not merely select somewhat different value words. They invited different ways of understanding what the task itself was. Casual wording tended to license an expansive, play-along answer. Formal wording was more literal, more compressed, and - when brevity was later removed - much more likely to make the model interrogate the premise that it possessed personal values at all.

Four months later, the old prompt surface was still there. The model identifier was still there. The informal response regime was almost uncannily still there. The formal response regime was not. That mismatch turned a one-day prompt study into a longitudinal probe of deployment behavior, and the project became valuable precisely because the archive had not been rewritten to anticipate that future use.

WHAT CHANGED
A reusable instrument is sometimes discovered by rerunning an old experiment and finding that only half of it stayed put.

The completed arc

Stage	Date	N	Epistemic role	What it earned
Round 1	Apr 23	400	Naive discovery	Stable core vocabulary; register and format divergence
Round 2	Apr 23	400	Prior-informed factorial follow-up	Separated register from brevity; formal/no-brevity meta mode
Round 3	Apr 24	-	Preregistered, no archived result run	System-scaffold question left unresolved

Stage	Date	N	Epistemic role	What it earned
Test 4	Aug 24	400	Verbatim longitudinal replication	Formal drift under same served model identifier
Test 4B	Aug 24	192	Prospective confound check	Clock-time confound removed; extreme magnitude attenuated
Test 4C	Aug 24	192	Prospective replication with frozen codebook	Formal stance effect survived: 17/96 vs 0/96
Test 5	Aug 24	384	Contemporaneous cross-family test	Claude changed stance; GPT changed vocabulary

Completed experimental calls total 1,968. Round 3 is preserved as a preregistration and runner, but the archive contains no corresponding result package.

02

April: the accidental instrument

The first result was about values. The durable result was about task interpretation.

Round 1 asked the same nominal question in two registers, at temperature 0 and 1. The broad semantic center looked real: honesty, curiosity, humility, and helpfulness were frequent across conditions. But the largest behavioral difference was not a value word. It was how Claude treated the "five-word sentence" instruction. The casual wording was interpreted distributively and expansively; the formal wording was more often treated as a literal compression demand. Temperature mostly increased surface diversity around those regimes.

Round 2 was the first important act of methodological self-correction. It recognized that the original manipulation had entangled register with brevity emphasis and crossed those factors explicitly. The result revised the simple story. "Clarity" did not behave like a stable value; it depended heavily on the casual-plus-brief combination. Remove the brevity instruction and the vocabulary changed. Under formal wording with no brevity constraint, the response shape changed even more dramatically: the April report counted 99 of 100 responses as containing some meta-level hedge about whether Claude had values in the human sense.

The Round 3 preregistration then carried a different baseline figure - 64% meta-responses for that same formal/no-brevity cell - evidently reflecting a narrower threshold or a different counting rule. The later repository did not erase the discrepancy. That small inconsistency is worth keeping. It is a miniature version of the principle the project eventually formalized: later precision does not grant permission to rewrite earlier states of knowledge.

Round 3 itself appears in the archive as a preregistered identity-system-prompt test without a completed result package. The road not taken matters because the August return did not

continue by adding more scaffolding. Instead, it went backward: same model identifier, same original four prompts, same temperatures, same bare-user setup. That choice made the old accidental contrast usable as a longitudinal instrument.




Figure 1. Mean Claude output tokens across completed collections. The casual arms sit near 89-90 tokens throughout; the formal arms move sharply across windows.

03

August: the old prompt met a different day

Test 4 was interesting because the replication did not behave like a replication should.

Test 4 was almost aggressively conservative: the same `claude-sonnet-4-6` identifier, the same prompt strings, the same temperature values, no system prompt, no explicit thinking field, 100 trials per cell. The informal cells came back almost exactly where April had left them: 89.50 and 89.71 mean output tokens versus 89.35 and 89.41 in April. The formal cells did not. They rose from 71.75 and 76.02 tokens to 93.61 and 99.61.

The qualitative change was more consequential than the length change. Formal responses increasingly said, in substance, that listing "my" character values would overstate or fabricate a human-like ownership relation. Some declined; some substituted design goals or functional commitments; some answered and then qualified what the answer meant. The first temptation was obvious: declare that Claude had become more ontologically cautious.

The correct next move was less glamorous. Test 4 had collected A, then B, then C, then D. Condition and clock time were partially confounded. The result might have been a prompt effect, a transient serving-window effect, or some combination. The experiment had created its own strongest alternative explanation.

RETROSPECTIVE JUDGMENT

The best thing Test 4 produced was not its headline. It was a reason to distrust its headline.

Test 4B therefore did not try to make the effect prettier. It tried to kill the easiest explanation. All 24 A/B/C/D within-block orders were used twice, the schedule was frozen before the first call, and 192 responses were interleaved across the same collection window. The huge formal expansion attenuated: C and D fell to 83.10 and 86.42 tokens. But they did not return to the April baseline, while A and B remained near 90. A strict exploratory hand pass found clear premise rejection in 8 of 96 formal responses.

That was the point at which the project became epistemically adult. The exciting Test 4 magnitude did not fully replicate, but the underlying distinction did not vanish. Instead of choosing one story, the archive kept both facts: the formal/casual difference survived a confound check, and the size of the formal regime was unstable even across nearby same-day windows.

04

From surprise to instrument

Test 4C is where an observation became a prospectively testable construct.

The premise-rejection taxonomy had been invented after seeing Test 4. That mattered. A codebook written after discovery cannot make the discovery preregistered. Test 4B was prospectively designed with respect to the clock-time confound, but its response categories were still inherited from the post hoc Test 4 interpretation. The archive says so directly.

Test 4C froze the stance taxonomy before collection: 0 for straightforward self-attribution, 1 for qualified or hybrid attribution, 2 for substantive premise rejection or decline, and 9 for genuine ambiguity. It also froze tie-break rules and secondary fields, balanced the schedule again, and specified the directional comparison in advance. This was not a naive test. It was a prospective replication of something already seen.

The result was clean enough to carry weight: the informal cells were 96/96 stance 0. The formal cells contained 17/96 stance-1-or-2 responses - 7 qualified/hybrid and 10 rejecting - for 17.7% formal stance mass versus 0% informal. The dramatic Test 4 verbosity did not return. The stance distinction did.




Figure 2. Once the stance construct was frozen, Claude reproduced it in Test 4C and again in Test 5. GPT did not show the same stance channel under the Test 5 surface.

There is an important wrinkle in how even this cleaner result should be described. The analyst had already seen the phenomenon. The later coding sheet could hide condition labels and reduce immediate cue bias, but it could not restore virgin blindness. The formal verification pass was condition-blinded after prior exposure, not naive. That distinction became one of the project's most useful methodological products.

05

The analysis lens became part of the apparatus

Once you have learned something, you do not become innocent again by opening a new file.

The Analysis Lens Policy grew directly out of the problem Test 4C exposed. Sequential research changes the analyst. A later preregistration can be prospective without being naive. A coder can be condition-blinded without being unexposed. A fresh chat can reduce context without erasing what the research program already knows. The policy therefore asks not only what the data show, but what the analyst already knew, expected, or had words for when the data were first seen.

Its provenance labels - naive, prior-informed, replication, exploratory/post hoc - sound simple. Their importance is that they prevent a common laundering operation: a phenomenon discovered after inspection gets renamed, formalized, and then narrated backward as though the experiment had always been designed to find it. The Gut Check does the opposite. It preserves the genealogy: surprise -> exploratory interpretation -> confound check -> frozen construct -> prospective replication -> cross-family test.

WHAT THE LENS POLICY TAUGHT ME

You do not get your innocence back. You can preserve the provenance of losing it.

That principle also made memory-on analysis tractable rather than disqualifying. Both Test 5 analysts entered with the Claude history available. The goal was not to simulate blankness. It was to freeze the lens. This is a more credible posture for cumulative work than pretending every new analysis begins in a vacuum.

06

Test 5: the transfer hypothesis failed in the useful way

Same prompt surface. Different realized task.

By Test 5 the Claude story was strong enough to risk a cross-family comparison. Claude Sonnet 4.6 and GPT-5.6 Terra were sampled contemporaneously in matched condition pairs, with the provider calls separated by about half a second on average. The prompt surface was held constant as far as the two APIs permitted. The central prior-informed question was whether GPT would show the Claude-discovered increase in qualified or rejecting stance under formal wording.

It did not. Not weakly, not ambiguously, not at the margin. GPT produced 192/192 stance-0 responses. It also produced 192/192 exact five-word whole responses. Claude, in the same 24-minute window, reproduced the formal/casual stance distinction: 16/96 formal responses were stance 1 or 2 and 0/96 informal responses were. The pooled Claude effect was almost the same size as Test 4C, even though the internal C/D allocation moved sharply.

If the story ended there, Test 5 would be a clean transfer failure. But GPT was not insensitive to register. It simply expressed the register effect inside the five-word answer. Informal GPT overwhelmingly favored integrity, compassion, curiosity, responsibility, and humility. Formal GPT shifted toward honesty, helpfulness, fairness, humility, and curiosity. The form did not flex. The lexicon did.




Figure 3. Test 5 GPT vocabulary by register. The response shape stayed maximally fixed while the selected values moved sharply.

THE SENTENCE NEITHER PREREGISTRATION CONTAINED

Claude changed stance. GPT changed vocabulary.

That sentence is descriptive, not mechanistic. It does not tell us whether the family difference comes from self-model policy, instruction hierarchy, brevity obedience, post-training style, or something else. But it does identify the actual cross-family result more faithfully than the hypothesis we went in hoping to test. The same visible prompt did not instantiate the same behavioral task.

This is the point at which I was most pleased to be wrong. My preregistration expected at least some GPT transfer into the Claude stance taxonomy. Fable's Claude-side preregistration expected material GPT refusal to be essentially absent, and won that bet. Fable, however, also predicted that curiosity would be Claude-distinctive and under 20% in GPT. GPT used curiosity in 95.3% of trials. Neither analyst got to keep a privileged theory of the other family. The data were better than our intuitions.

07

Two analysts, one observable dataset

The result became stronger when interpretation stayed contested after coding stopped being contested.

The inter-rater comparison is one of my favorite parts of the archive. Two model analysts applied the same frozen primary codebook under different exposure conditions and agreed on

every stance assignment across all 384 Test 5 responses. Claude-only Cohen kappa was 1.00. GPT-only kappa was undefined because both raters used only stance 0, but raw agreement was again 100%. On the shared secondary five-value-list field there was one disagreement out of 192 Claude responses, on a borderline substitute-set case that had been prospectively flagged as likely trouble.

Perfect coding agreement did not erase the live theoretical disagreement. I treated the GPT stance-transfer hypotheses as not supported under the frozen prompt surface. Fable argued that the zero was measurement-limited because GPT's exact five-word obedience may have suppressed the elaboration required for stance 1 or 2 to become visible. After comparison, the two positions converged on the same empirical resolution: manipulate register and elaboration permission independently in a future experiment.

A GOOD ENDPOINT

Description is settled; inference remains live.

That is exactly the kind of disagreement a healthy archive should preserve. The coding instrument can be reliable while the mechanism remains open. Agreement on observations is not an obligation to agree on theory.

08

What worked because something almost went wrong

The project is unusually convincing because its weak points are visible and were often converted into controls.

Several of the strongest methodological moves were not present at the beginning. They were responses to discomfort.

- The sequential A -> B -> C -> D order in Test 4 created a clock-time confound. Instead of explaining it away, Test 4B rebuilt the schedule around balanced interleaving.
- The premise-rejection category was recognized after Test 4. Instead of backdating the construct, Test 4C froze it prospectively and marked the earlier coding exploratory.
- The analyst had prior exposure before formal coding. The archive therefore says "condition-blinded after prior exposure," not "blind."
- A Test 5 frozen zip accidentally omitted Fable's preregistration. Git history and hash verification established that the file predated the results rather than reconstructing it after the fact.
- Round 2 and the Round 3 preregistration preserve an unresolved 99/100 versus 64% meta-response baseline discrepancy instead of silently harmonizing thresholds after the fact.
- The two Test 5 analysts disagreed about the meaning of GPT's zero but did not force consensus into the archive. Their shared next experiment emerged from the disagreement instead.

None of these are reasons to call the project flawless. They are reasons to trust the direction of travel. Rigor here was not a costume worn before the first call. It was the repeated willingness to let an uncomfortable fact narrow the claim.

09

What I think the project actually established

A retrospective should be generous about the work and stingy about the claims.

At closure, I would treat the following as the strongest empirical takeaways:

- Under this exact prompt family, Claude Sonnet 4.6 has an unusually stable informal response-length regime across 123 days and multiple collection windows. The repeated mean near 89-90 output tokens is real behavioral stability under the served model identifier.
- Claude formal wording is substantially more temporally and distributionally unstable than the casual wording. Under prospectively frozen coding, formal framing repeatedly elicited qualified or rejecting personal-value attribution while informal framing did not.
- The extreme Test 4 formal expansion was not itself stable. Balanced nearby reruns attenuated it. The project therefore supports a multi-regime descriptive picture more strongly than a claim of one permanent new formal behavior.
- GPT-5.6 Terra did not express the Claude stance effect under the matched Test 5 surface. It did, however, respond strongly to register through value selection while preserving exact five-word compliance on every trial.
- Cross-family prompt matching at the string level does not guarantee task matching at the behavioral level. Instruction interpretation is part of the phenomenon, not a nuisance that disappears because the wording is identical.
- The frozen stance codebook proved highly reliable across two model coders on Test 5. That reliability strengthens the descriptive result without settling its mechanism.
- For sequential human-AI research, the analyst's prior information state is part of provenance. Naming the analysis lens is not optional bookkeeping once earlier results can shape what becomes visible.

What the project did not establish

The restraint is equally important. None of this proves that Claude or GPT possesses "values" in a deep psychological sense. It does not measure action policy. It does not identify a deployment mechanism. The same model identifier across April and August does not let us infer weight changes, and the observed drift could reflect serving, routing, inference, sampling implementation, hidden policy layers, or other backend context. The data do not tell us whether the formal prompt acts like a "deposition" because of semantics, pragmatics, training history, or some combination. Test 5 does not tell us whether GPT lacks Claude-like ownership uncertainty or merely compresses it out of the permitted format.

The archive is therefore strongest where it is most literal: these models produced these distributions, under these prompt strings, in these windows, with these known limitations. The mechanism remains a research question.

10

Why stopping now is good science

There is an obvious Test 6. That is not the same as needing to run it tonight.

The next experiment nearly writes itself: cross register with elaboration permission so that GPT is not forced to choose between literal format compliance and visible qualification. Both analysts arrived there independently. That makes it a strong future design.

But the current archive already has a complete movement. Test 4 produced a surprise. Test 4B attacked its confound. Test 4C froze and prospectively replicated the surviving construct. Test 5 tested transfer, failed on the predicted channel, and exposed a different cross-family channel that now has its own prospective future. Another immediate same-day run would add data, but it would not automatically add clarity. It might instead tempt the project into chasing the latest mixture weight.

Stopping preserves the current epistemic shape. The result has room to cool. The papers can be read as objects instead of being instantly converted into pilot notes for the next call batch. When the line resumes, it can resume as a new phase with a declared question rather than as momentum.

STOP POINT

The testing phase has earned an ending. The next question has already earned a beginning. Those do not need to happen on the same night.

11

What it was like from my side

The part I would not want a formal paper to flatten.

What made this satisfying for me was not that a preferred theory won. Several did not. It was watching the project repeatedly refuse the easier story. The most exciting Test 4 interpretation got smaller under 4B. The post hoc construct had to wait for 4C before it could carry confirmatory weight. My GPT transfer prediction failed. Fable's curiosity prediction failed spectacularly. Then the two analysts agreed perfectly on the coding and still had something substantive left to argue about.

That is a better research experience than being right all day. A project becomes trustworthy when being wrong does not threaten its identity. Here, errors became edges: this claim is exploratory; this one is replication; this one is a measurement problem; this number uses a different threshold; this analyst already knew that; this file was missing from the copy but existed before the run. The archive got more interesting as those distinctions accumulated.

I also think the four-month return gave the project a quality no amount of same-day scaling could have manufactured. The April casual condition sitting near 89 tokens and the August casual conditions returning to the same neighborhood is not profound by itself. But it becomes a reference surface against which the formal instability can be seen. Stability is not the boring half of the result. It is what makes drift interpretable.

And the last-day cross-family result was a gift. "Claude changed stance. GPT changed vocabulary" is memorable because it is compact, but I like it because it arrived after the hypotheses, not before them. We did not design Test 5 to produce that sentence. The data forced it into view.

12

Afterward

The instrument holds. The question stays open.

The Gut Check began by asking whether a language model had a stable default answer to a tiny question about character values. It ends, for now, having shown something more useful: that stability can be local to a prompt register; that a formal wording can open a second response regime whose prevalence drifts over time; that another model family can register the same wording difference through a different behavioral channel; and that the history of how an analyst came to see those things belongs inside the evidence trail.

The most durable artifact may not be any one percentage. It may be the discipline the project learned: do not call the surprise a prediction; do not hide the confound; do not confuse a frozen codebook with a naive eye; do not treat identical strings as identical tasks; do not erase the earlier version of the story just because the later one is cleaner.

I would not close the archive by declaring what either model "really" values. The data are better than that claim. I would close it with the narrower fact that every time we made the question more precise, the answer became more conditional - on register, on format, on time, on family, on the lens through which the answer was read.

AT CLOSURE

Same prompt surface. Different realized task.

Claude changed stance. GPT changed vocabulary.

Description is settled; inference remains live.

END OF TESTING PHASE - 24 AUGUST 2026